

lqr

October 4, 2019

```
In [62]: from IPython import display
import matplotlib.pyplot as plt
from matplotlib.lines import Line2D
from gym.envs.mujoco import *
from envs.hopper_env import HopperModEnv
from envs.cheetah_env import CheetahModEnv
import numpy as np
import copy
import gym
from scipy.io import loadmat
from scipy.io import savemat
import moviepy.editor as mpy
from simulators import *
from rot_utils import *
import seaborn as sns
sns.set_style('darkgrid')
import warnings
warnings.filterwarnings('ignore')
```

(a) LQR for Linear Systems Let's start with a linear system:

```
In [63]: A = np.array([[0.0481, -0.5049, 0.0299, 2.6544, 1.0608],
                        [2.3846, -0.2312, -0.1260, -0.7945, 0.5279],
                        [1.4019, -0.6394, -0.1401, 0.5484, 0.1624],
                        [-0.0254, 0.4595, -0.0862, 2.1750, 1.1012],
                        [0.5172, 0.5060, 1.6579, -0.9407, -1.4441]])
B = np.array([[-0.7789, -1.2076],
              [0.4299, -1.6041],
              [0.2006, -1.7395],
              [0.8302, 0.2295],
              [-1.8465, 1.2780]])
dx = A.shape[0]
du = B.shape[1]
```

Now verify the system is controllable.

To do so we need to check if $\text{rank}([B \ AB \ A^2B \ A^3B \ A^4B]) = 5$. (generally we'd go up to $A^{(n-1)}B$, and verify $\text{rank} == n$)

```
In [64]: # verify the above statement
lst = [B]
"""YOUR CODE HERE"""
An = np.identity(dx)
for i in range(1, dx):
    An = An.dot(A)
    AnB = An.dot(B)
    lst.append(AnB)

"""YOUR CODE ENDS HERE"""
np.linalg.matrix_rank(np.hstack(lst))
```

Out[64]: 5

Recall the following optimal control problem:

$$\min_{x,u} \sum_{t=0}^{T-1} (x_t' Q x_t + u_t' R u_t) + x_T' Q_{final} x_T \text{ s.t. } x_{t+1} = A x_t + B u_t$$

For T going to infinity, we'll have that the Value Function and the Feedback Controller at time $t = 0$ reach a steady-state — the optimal value function and feedback controller for infinitely many time-steps to-go

So let's run the Value Iteration Solution to the LQR control problem until it has converged, and then use that infinite horizon optimal feedback controller to stabilize our system at 0.

```
In [65]: # implement the infinite horizon optimal feedback controller
def lqr_infinite_horizon(A, B, Q, R):
    """
find the infinite horizon K and P through running LQR back-ups
until l2-norm(K_new - K_curr, 2) <= 1e-4
return: K, P
"""

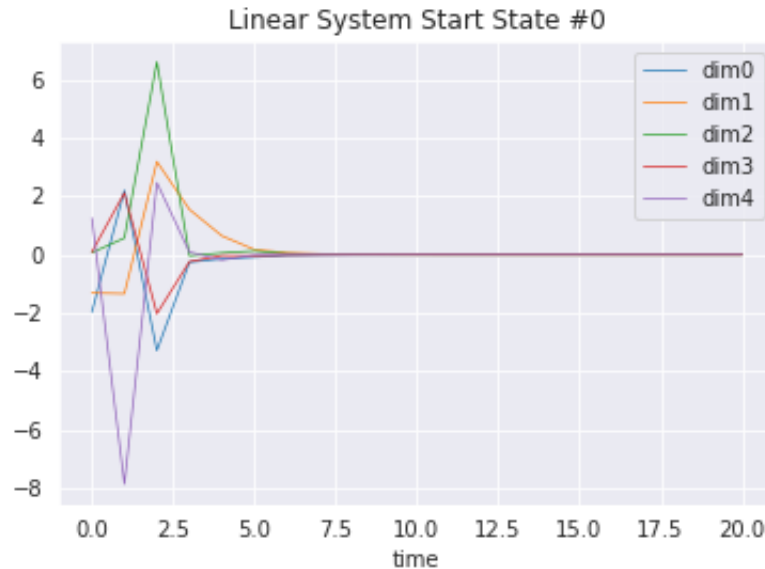
    dx, du = A.shape[0], B.shape[1]
    P, K_current = np.eye(dx), np.zeros((du, dx))

    """YOUR CODE HERE"""
    while True:
        K_new = - np.linalg.inv(R + B.T @ P @ B) @ B.T @ P @ A
        P = Q + K_new.T @ R @ K_new + (A + B @ K_new).T @ P @ (A + B @ K_new)
        if np.linalg.norm(K_new - K_current, 2) <= 1e-4:
            break
        else:
            K_current = K_new

    """YOUR CODE ENDS HERE"""
    return K_new, P

In [66]: # problem has been defined, let's solve it:
Q, R = np.eye(dx), np.eye(du)
K_inf, P_inf = lqr_infinite_horizon(A, B, Q, R)
```

Now let's simulate and see what happens for a few different starting states.
 Here's what a reference plot looks like for the *first starting state with no noise* in state dynamics



so you may compare:

```
In [67]: # fill in the simulation to use your controller, K_inf, at each timestep then run the c
def simulate(A, B, K_inf, n_starting_states, T, noise=None):
    for s in np.arange(n_starting_states):
        x, u = np.zeros((K_inf.shape[1], T+1)), np.zeros((K_inf.shape[0], T+1))
        x[:,0] = starting_states[:,s]
        for t in np.arange(T):
            """YOUR CODE HERE"""
            u[:,t] = K_inf @ x[:,t]
            """YOUR CODE ENDS HERE"""
            x[:,t+1] = A @ x[:,t] + B @ u[:,t]
            if noise is not None:
                x[:,t+1] += noise[:,t]
        plt.plot(x.T, linewidth=.7)
        plt.xlabel('time')
        plt.title("Noisy Linear System Start State #{}".format(s)) if noise is not None
        plt.legend(["dim"+str(i) for i in range(len(x))])
        plt.show()

starting_states = np.array([[-1.9613, 1.9277, -0.2442],
                             [-1.3127, -0.2406, -0.0260],
                             [0.0698, -0.5860, -0.7522],
                             [0.0935, -0.1524, -0.9680],
                             [1.2494, 0.5397, -0.5146]])

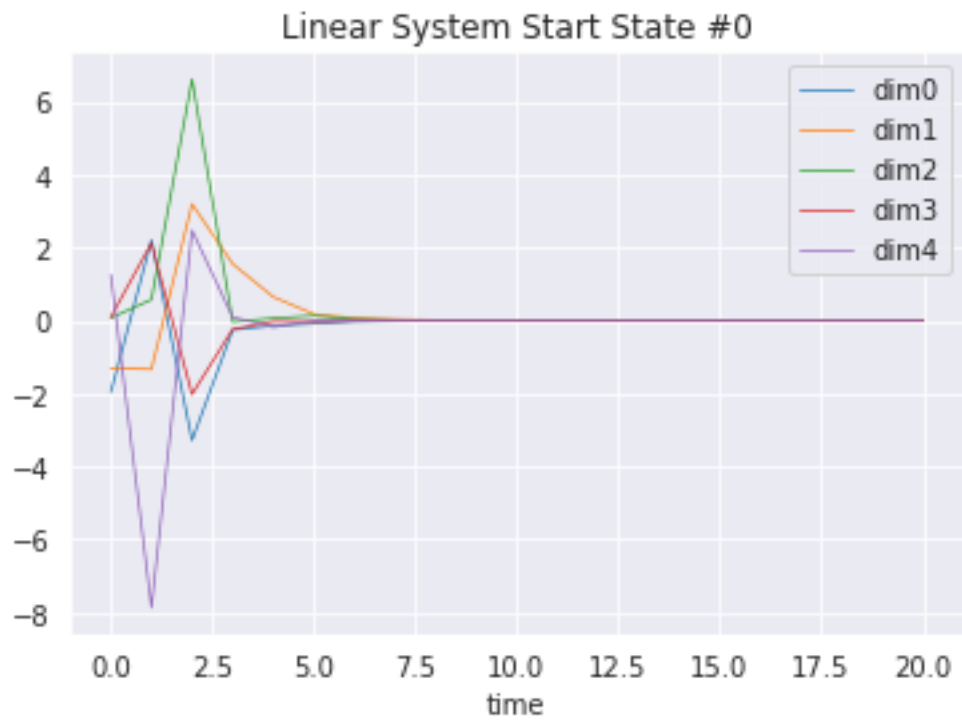
n_starting_states = starting_states.shape[1]
T = 20 # simulating for 20 steps
simulate(A, B, K_inf, n_starting_states, T)
```

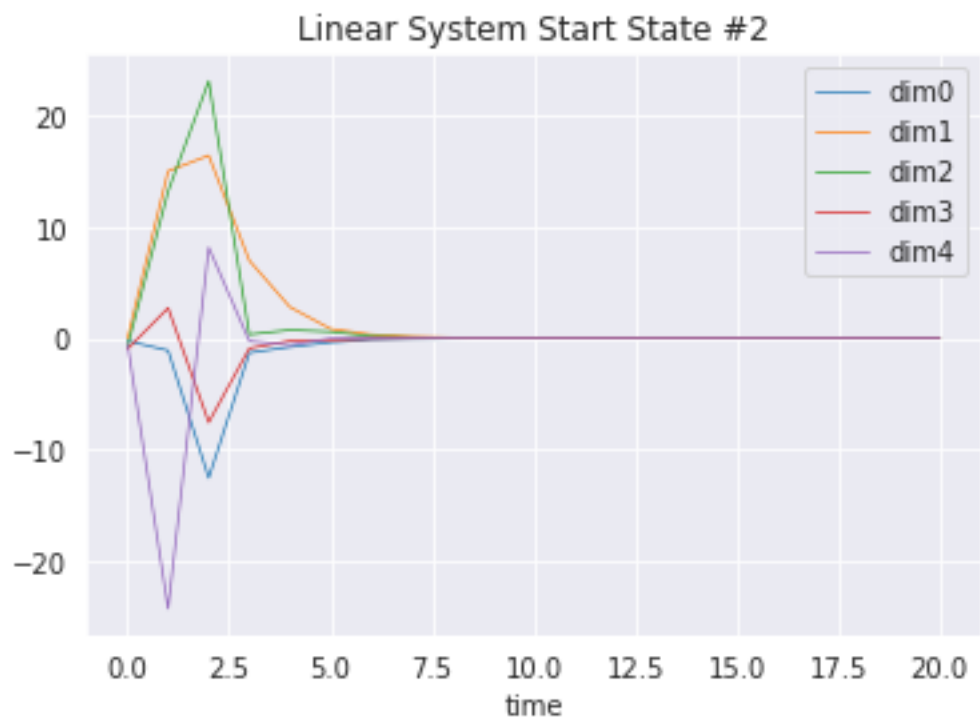
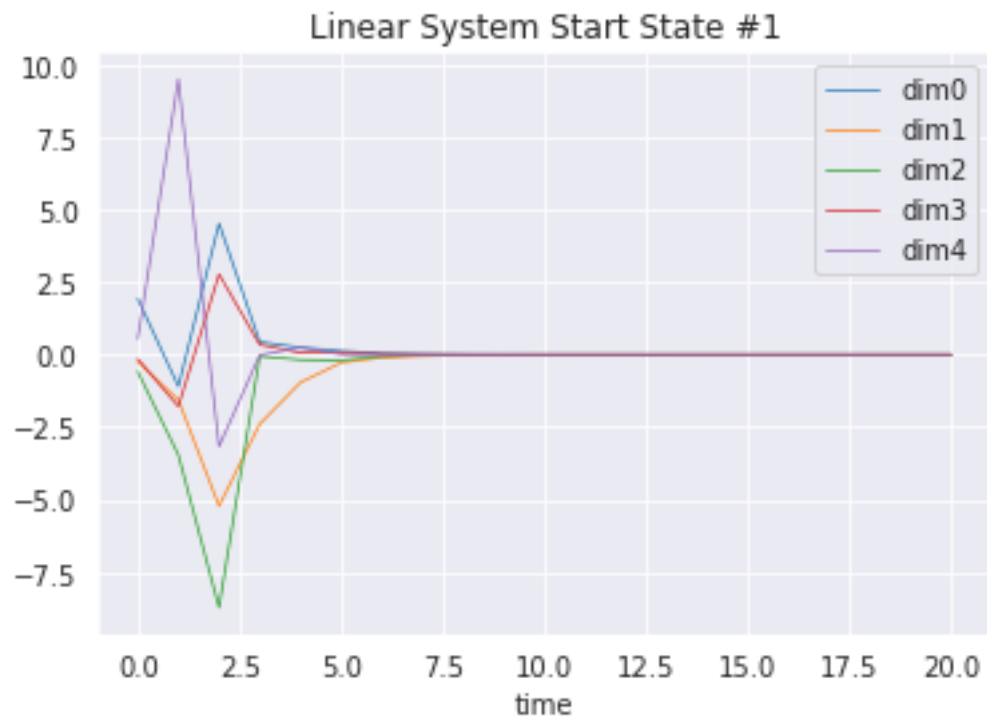
and in the presence of noise:

noise_id = "p_a_w"

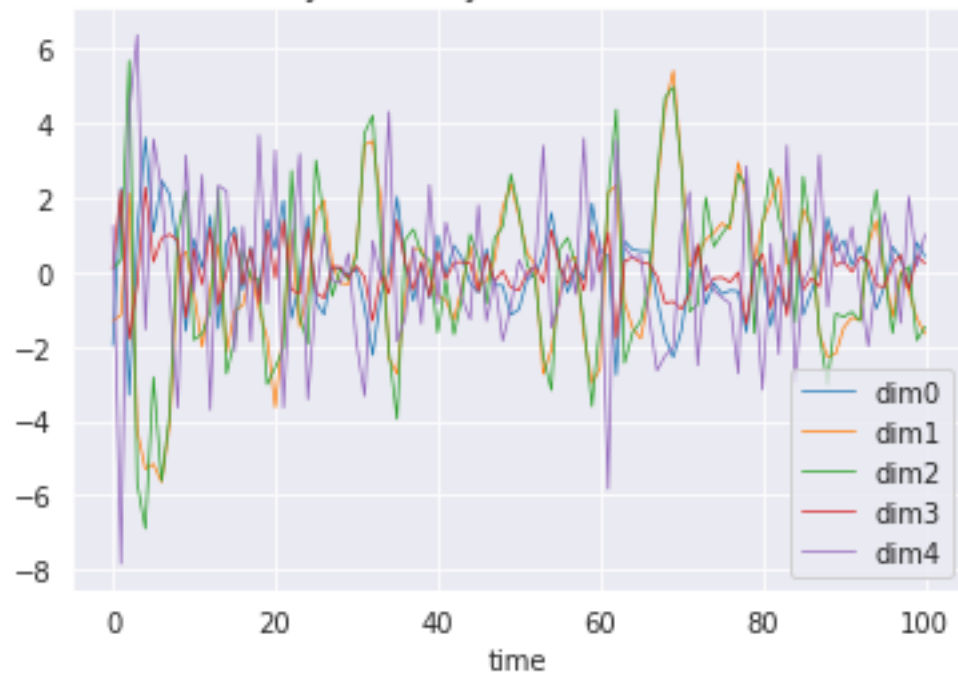
T = 100 *# simulating for 100 steps*

simulate(A, B, K_inf, n_starting_states, T, noise=loadmat("mats/"+noise_id+".mat")[noise_id])

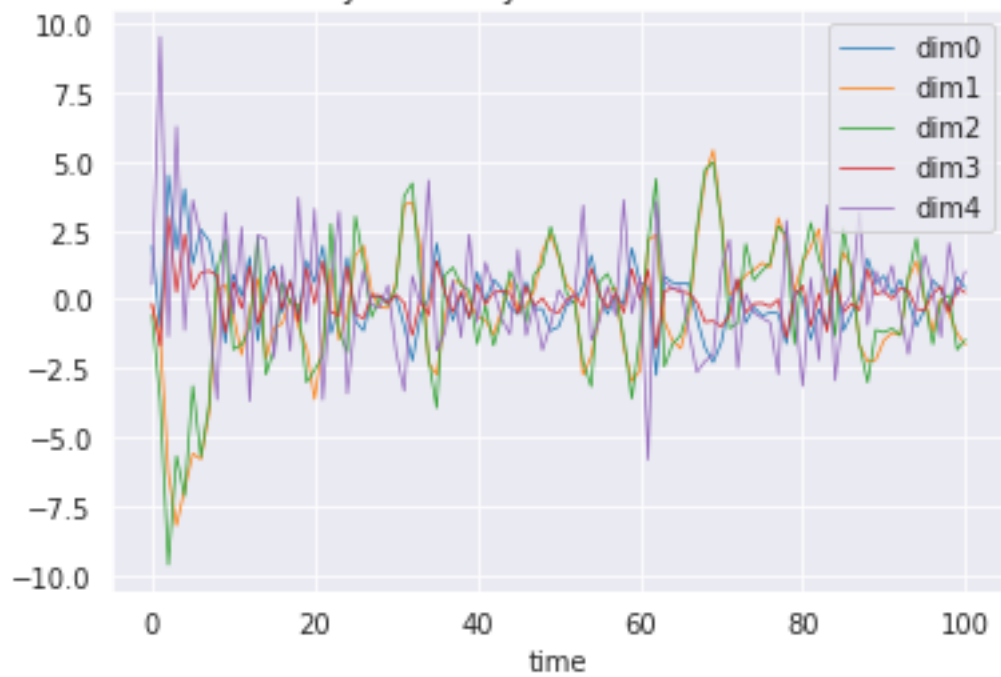


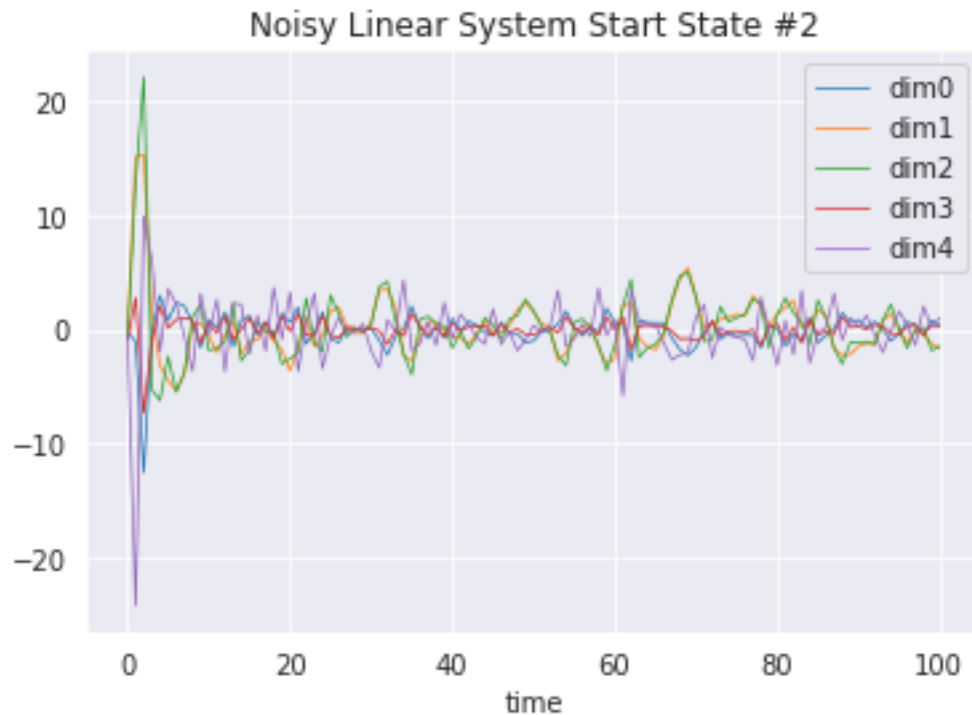


Noisy Linear System Start State #0



Noisy Linear System Start State #1





(b) LQR-based Stabilization for Nonlinear Systems Now let's consider nonlinear systems. Linearize around one point and design an infinite horizon controller for the resulting system.

In [72]: *# implement linearization about a point*

```
def linearize_dynamics(f, x_ref, u_ref, dt, my_eps, x_ref_tplus1=None):
    """
    f : dynamics simulator
    my_eps : delta for forward and backward differences you'll need
    NOTE: please use centered finite differences!

     $x(:,t+1) - x_{ref} \text{ approximately } = A*(x(:,t)-x_{ref}) + B*(u(:,t) - u_{ref}) + c$ 
    If we pick  $x_{ref}$  and  $u_{ref}$  to constitute a fixed point, then  $c == 0$ 

    For part (b), you do not need to use the optional argument (nor c).
    For part (d), you'll have to revisit and modify this function
        --at this point, you'll want to use the optional argument and the resulting c.

    return: A, B, c
    """

    if x_ref_tplus1 is not None:
        x_ref_next = x_ref_tplus1
    else:
```

```

x_ref_next = x_ref

dx, du = x_ref.shape[0], u_ref.shape[0]
A, B = np.zeros((dx, dx)), np.zeros((dx, du))

"""YOUR CODE HERE"""
for j in range(dx):
    xj_forward = x_ref.copy()
    xj_forward[j] = xj_forward[j] + my_eps / 2
    xj_backward = x_ref.copy()
    xj_backward[j] = xj_backward[j] - my_eps / 2
    A[:, j] = (f(xj_forward, u_ref, dt) - f(xj_backward, u_ref, dt)) / (my_eps)

for j in range(du):
    uj_forward = u_ref.copy().astype(float)
    uj_forward[j] = uj_forward[j] + my_eps / 2
    uj_backward = u_ref.copy().astype(float)
    uj_backward[j] = uj_backward[j] - my_eps / 2
    B[:, j] = (f(x_ref, uj_forward, dt) - f(x_ref, uj_backward, dt)) / (my_eps)
"""YOUR CODE ENDS HERE"""

c = f(x_ref, u_ref, dt) - x_ref_next
if len(B.shape) == 1:
    return A, B.reshape(-1, 1), c
return A, B, c

```

In [73]: *# take an environment and find the infinite horizon controller for the linearized system*

```

def lqr_nonlinear(config):
    env = config['env']
    f = config['f']
    dt = 0.1 # we work with discrete time
    my_eps = 0.01 # finite difference for numerical differentiation

    # load in our reference points
    x_ref, u_ref = config['x_ref'], config['u_ref']

    # linearize
    A, B, c = linearize_dynamics(f, x_ref, u_ref, dt, my_eps)
    dx, du = A.shape[0], B.shape[1]
    Q, R = np.eye(dx), np.eye(du)*2

    # solve for the linearized system
    K_inf, P_inf = lqr_infinite_horizon(A, B, Q, R) # you implemented in part (a)

    # recognize the simulation code from part (a)? modify it to use your controller at
    def simulate(K_inf, f, x_ref, u_ref, dt, n_starting_states, T, noise=None):
        for s in np.arange(n_starting_states):
            x, u = np.zeros((K_inf.shape[1], T+1)), np.zeros((K_inf.shape[0], T+1))

```



```

x[:,0] = starting_states[:,s]
for t in np.arange(T):
    """YOUR CODE HERE"""
    u[:,t] = u_ref + K_inf @ (x[:, t] - x_ref)
    """YOUR CODE ENDS HERE"""
    x[:,t+1] = f(x[:,t], u[:,t], dt)
    if "p_val" in config.keys():
        perturbation_values = config["p_val"]
        perturb = perturbation_values[t//(T//len(perturbation_values))]
        x[:,t+1] = f(x[:,t], u[:,t], dt, rollout=True, perturb=perturb)
    if env is not None:
        if t % 5 == 0:
            plt.clf()
            plt.axis('off')
            plt.grid(b=None)
            plt.imshow(env.render(mode='rgb_array', width=256, height=256))
            plt.title("Perturbation Magnitude {}".format(perturb))
            display.clear_output(wait=True)
            display.display(plt.gcf())

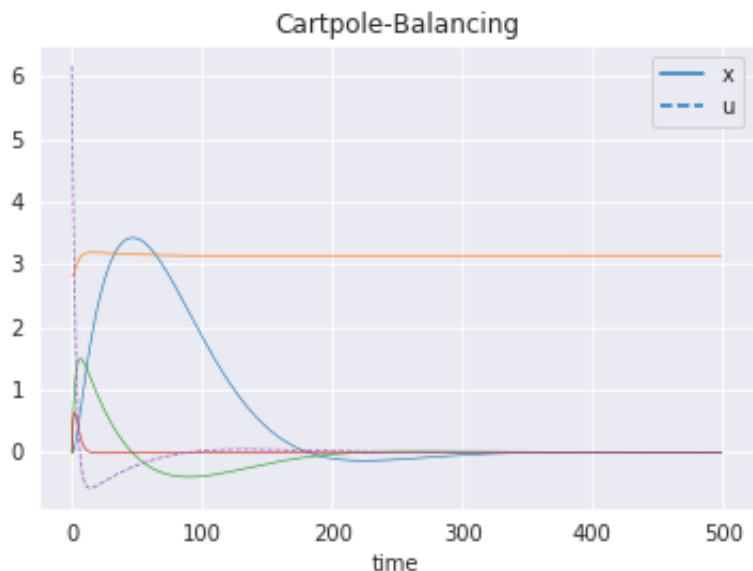
        if noise is not None:
            x[:,t+1] += noise[:,t]
    if env is not None:
        plt.clf()

plt.plot(x.T[:-1], linewidth=.6)
plt.plot(np.squeeze(u.T[:-1])/10.0, linewidth=.7, linestyle='--') # scaling
if 'legend' in config.keys():
    config['legend'].append('u')
    plt.legend(config['legend'])
else:
    legend_elements = [Line2D([0], [0], label='x'), Line2D([0], [0], linestyle='--', label='u')]
    plt.legend(handles=legend_elements)
plt.xlabel('time')
plt.title(config["exp_name"])
plt.show()

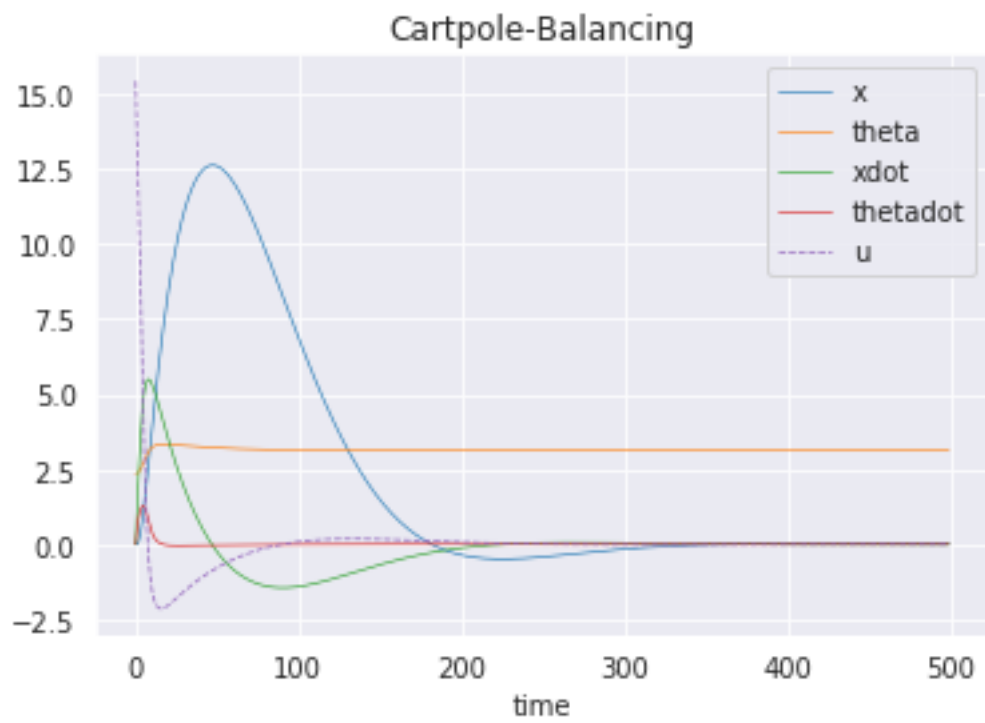
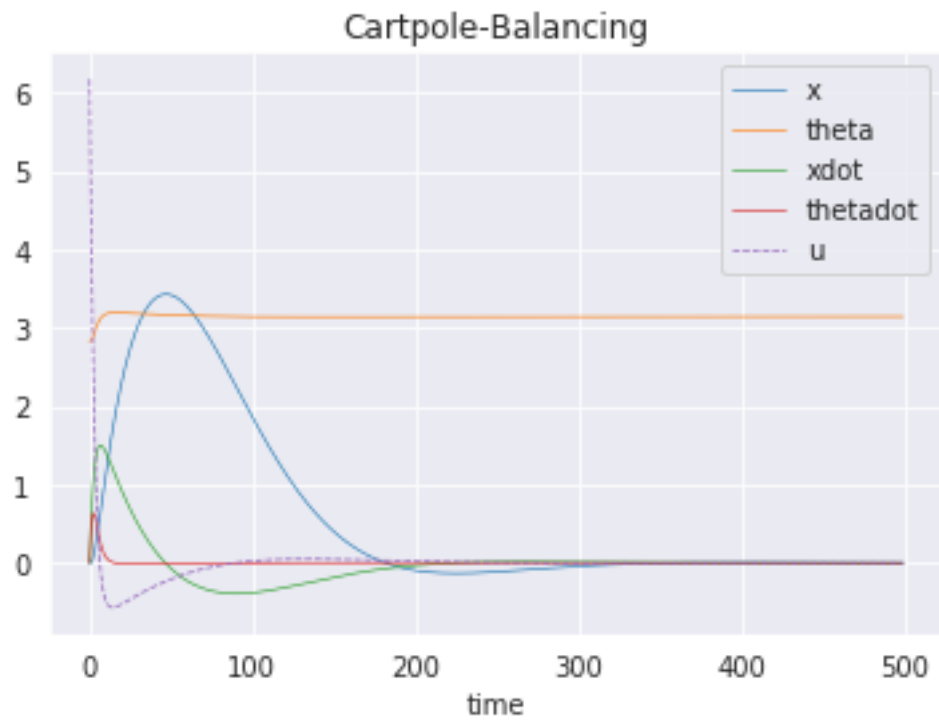
# now let's simulate and see what happens for a few different starting states
starting_states = config['ss']
n_starting_states = starting_states.shape[1]
T = config['steps'] # simulating for T steps
simulate(K_inf, f, x_ref, u_ref, dt, n_starting_states, T)
if 'noise' in config.keys():
    # and now in the presence of noise
    noise_id = config['noise']
    noise_loaded = loadmat("mats/"+noise_id+".mat")[noise_id]
    simulate(K_inf, f, x_ref, u_ref, dt, n_starting_states, noise_loaded.shape[1],

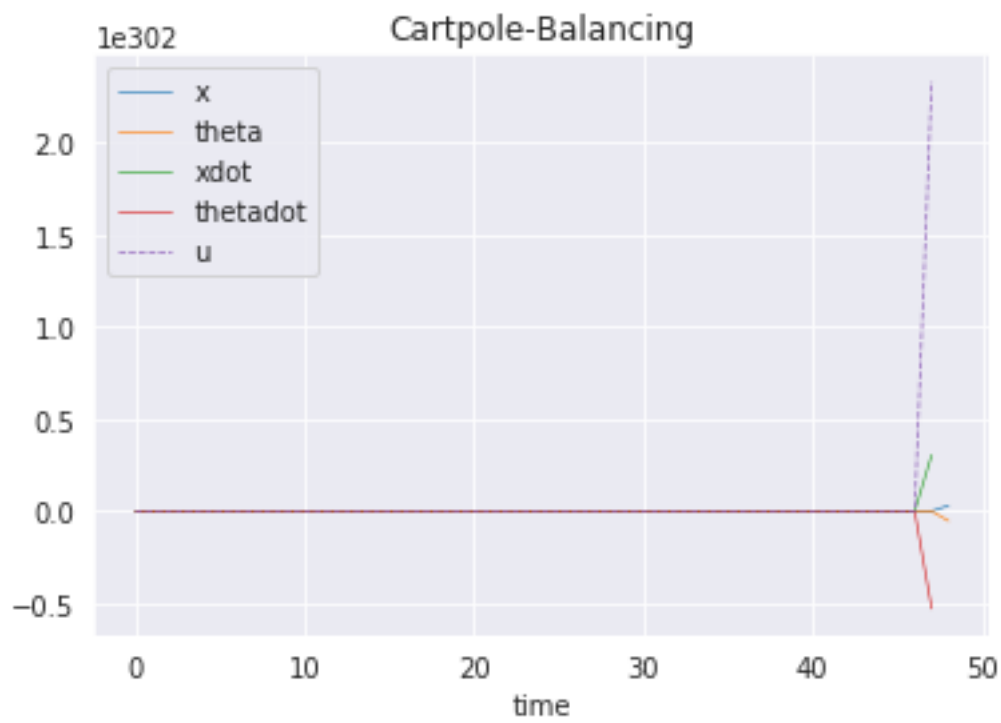
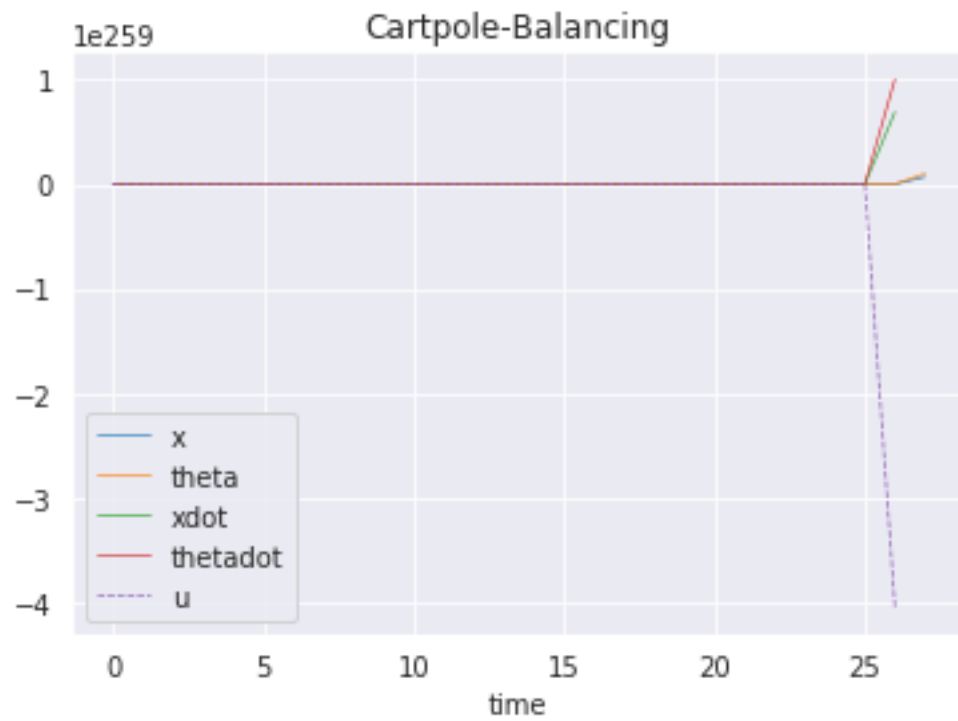
```

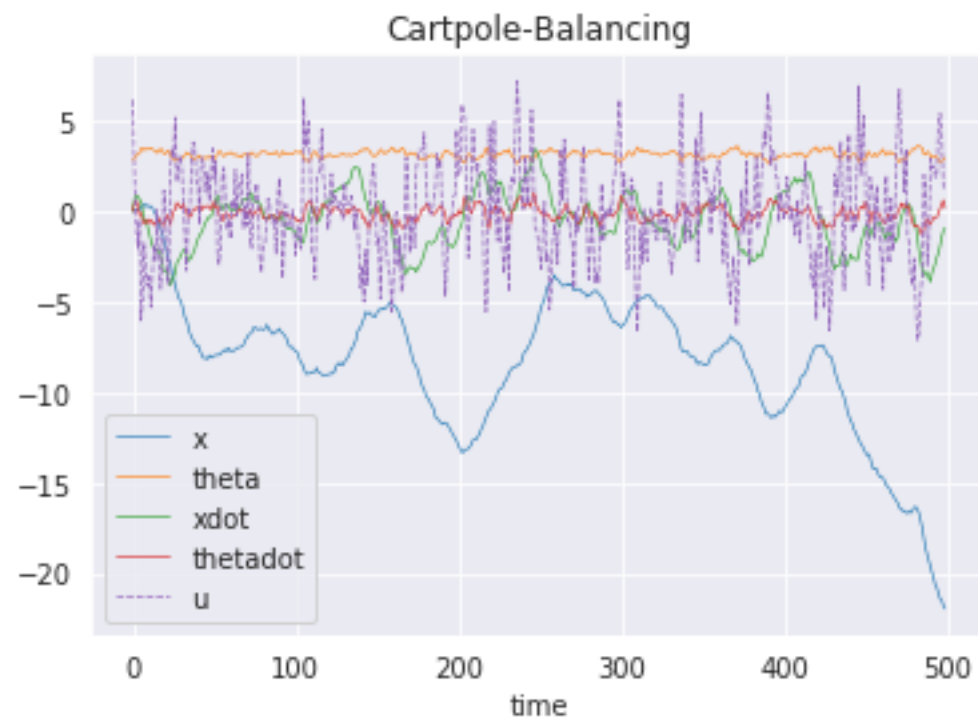
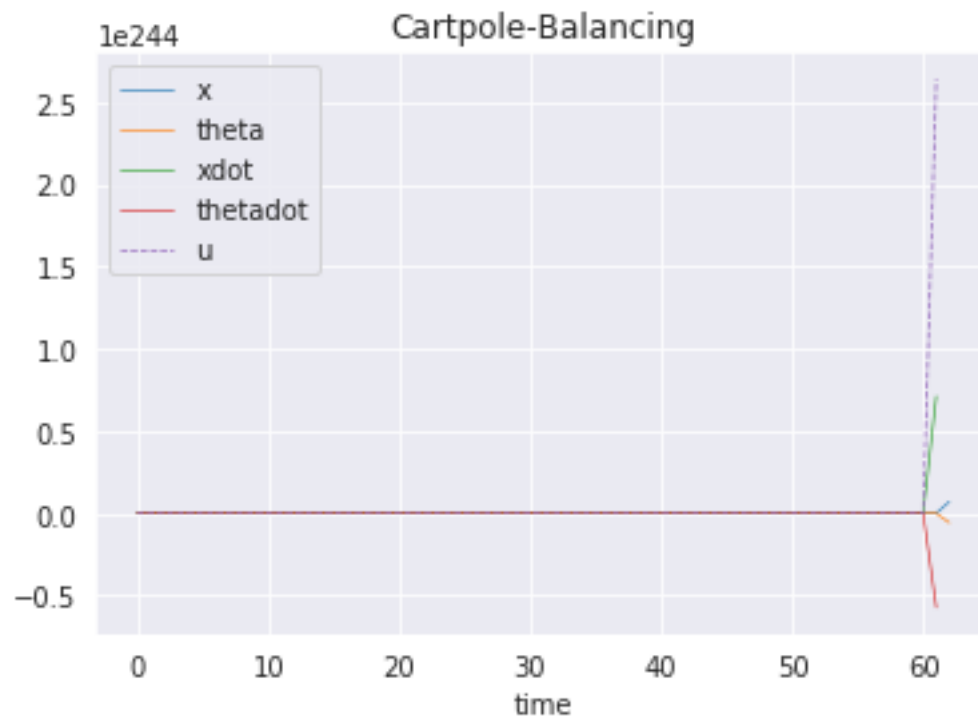
Cartpole-Balancing Here's what a reference plot looks like for the first starting state with no noise in state dynamics so you may compare:

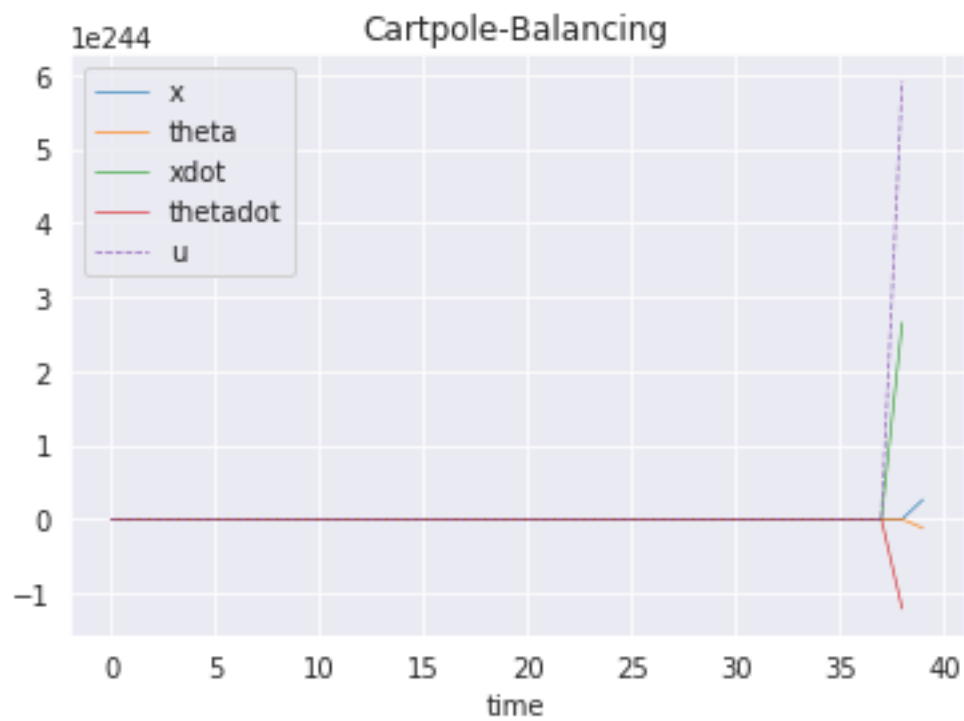
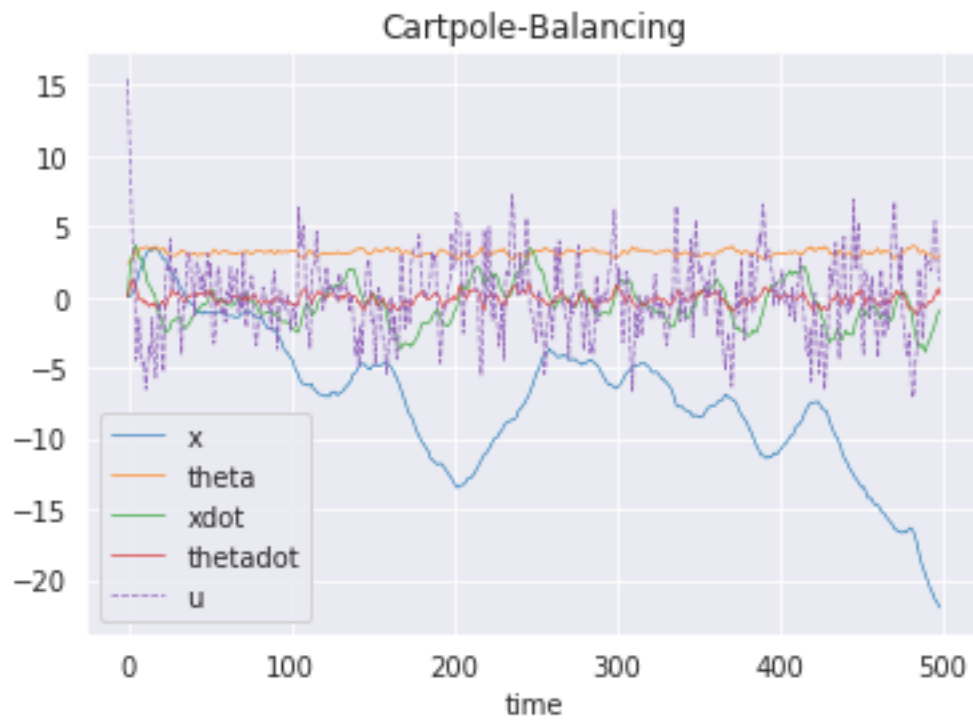


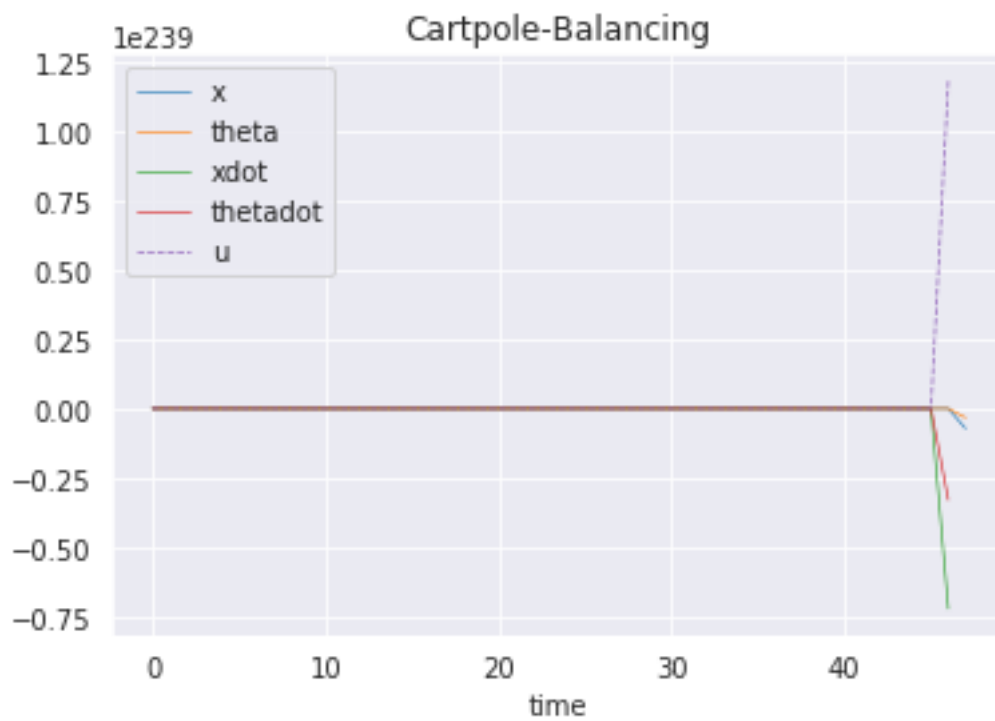
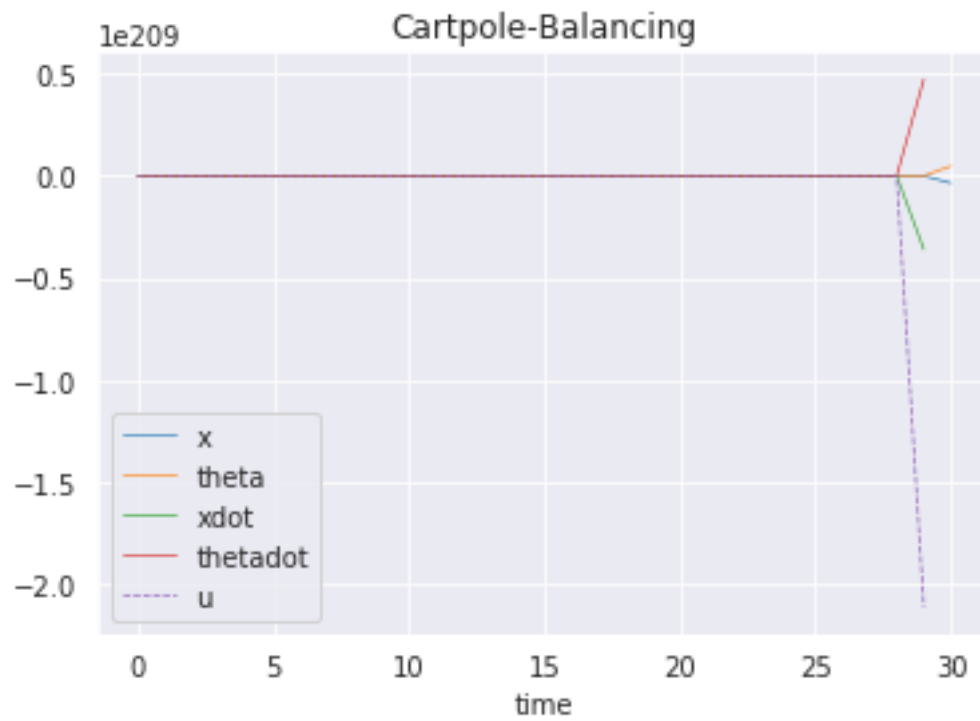
```
In [74]: # Find the infinite horizon controller for the linearized version of the cartpole balanc
cartpole_config = {
    'f': sim_cartpole,
    'exp_name': "Cartpole-Balancing",
    'env': None,
    'steps': 500,
    'x_ref': np.array([0, np.pi, 0, 0]),
    'u_ref': np.array([0]),
    'legend': ['x', 'theta', 'xdot', 'thetadot'],
    'ss': np.array([[0, 0, 0, 10, 50],
                    [9*np.pi/10, 3*np.pi/4, np.pi/2, 0, 0],
                    [0, 0, 0, 0, 0],
                    [0, 0, 0, 0, 0]]), #ss = starting states
    'noise': 'p_b_w',
}
lqr_nonlinear(cartpole_config)
```







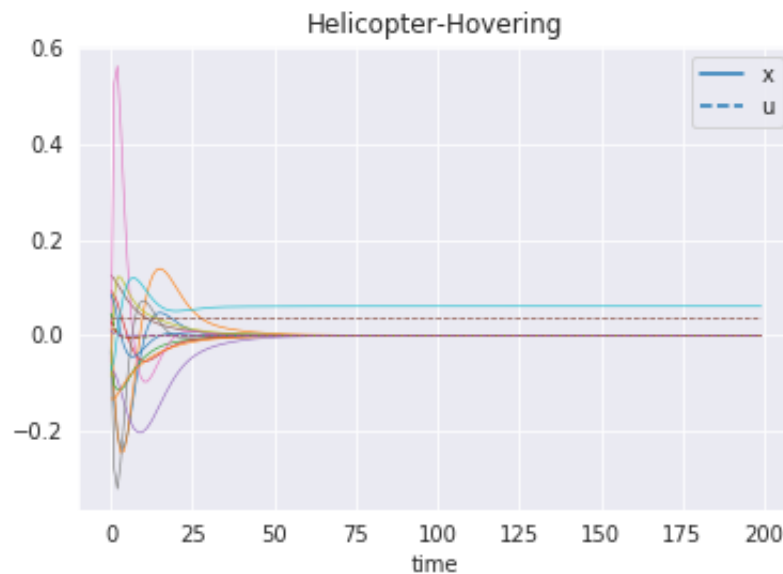




Question: Inspecting the generated plots for the system without noise, explain in 2-3 sentences how the different starting states affect the problem differently. In your answer, include an analysis of one starting state which causes failure and explain why this might happen.

Response: The starting states may be close to or far from the reference state. In the latter case, the linearization using Taylor's expansion drives the system far away from its original nonlinearity, leading to the inaccurate approximation. For instance, the last starting state is $[50, 0, 0, 0]$ which is very far from the reference state $[0, \pi, 0, 0]$, resulting in the diverging result.

Helicopter in Hover Now let's stabilize a helicopter in hover. We can use the *same* code written for the cartpole system to do so; simply run the cell below which defines the fixed point. Here's what a reference plot looks like for the *first starting state with no noise* in state dynamics so



you may compare:

```
In [10]: # Find the infinite horizon controller for the linearized version of the hovering copte
# Just run the cell below to generate plots using the code you wrote for cartpole!
x_ref, u_ref = np.zeros(12), np.zeros(4)
x_ref[9] = np.arcsin(3.0/(5*9.81))
u_ref[3] = 9.81*5*np.cos(x_ref[9])/137.5
heli_config = {
    'f': sim_heli,
    'env': None,
    'exp_name': "Helicopter-Hovering",
    'steps': 200,
    'x_ref': x_ref,
    'u_ref': u_ref,
    'ss': loadmat("mats/p_c_heli_starting_states.mat")["heli_starting_states"], #ss = s
    'noise': 'p_c_w',
}
lqr_nonlinear(heli_config)
```